

# AI Crop Monitoring & Agricultural Intelligence Platform – Full Implementation Guide

Generated by Grok

September 03, 2025

## Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
<b>2</b>	<b>Project Structure</b>	<b>3</b>
2.1	When & How You'll Use It	4
<b>3</b>	<b>Features &amp; Capabilities</b>	<b>4</b>
<b>4</b>	<b>Next-Level Differentiators</b>	<b>6</b>
<b>5</b>	<b>Governance, Compliance &amp; Reliability</b>	<b>6</b>
5.1	Licensing & Provenance	6
5.2	Annotation & Label QA	7
5.3	Data Validation	7
5.4	Model Governance & Cards	7
5.5	Drift Detection & Retraining	7
5.6	Security, Privacy & Compliance	8
5.7	Observability & SLOs	8
5.8	CI/CD & Testing Gates	8
5.9	Reproducibility & Backups	8
<b>6</b>	<b>Knowledge &amp; Skills Roadmap</b>	<b>8</b>
6.1	Mathematics & Statistics	8
6.2	Machine Learning	9
6.3	Deep Learning	9
6.4	MLOps	9
6.5	Edge AI	9
6.6	Compliance	9
<b>7</b>	<b>Frontend Details</b>	<b>9</b>
7.1	Technologies and Frameworks	9
7.2	Key Components and Pages	10
7.3	Code Structure	10
7.4	Example Code Snippet (Streamlit Upload)	10
7.5	Deployment and Testing	11

<b>8 Backend Architecture (Atlas Integrated)</b>	<b>11</b>
8.1 Technologies and Frameworks . . . . .	11
8.2 MongoDB Atlas Collections . . . . .	11
8.3 Use Cases . . . . .	11
8.4 Key Endpoints and Services . . . . .	11
8.5 Code Structure . . . . .	12
8.6 Example Code Snippet (FastAPI with MongoDB Atlas) . . . . .	12
8.7 Deployment and Testing . . . . .	12
<b>9 Edge &amp; Offline Devices</b>	<b>13</b>
9.1 Edge Models . . . . .	13
9.2 Caching . . . . .	13
9.3 Offline Recommendations . . . . .	13
9.4 Workflow Example . . . . .	13
<b>10 Observability &amp; CI/CD</b>	<b>13</b>
10.1 Monitoring . . . . .	13
10.2 Model Tracking . . . . .	13
10.3 CI/CD . . . . .	14
10.4 Alerts & Logs . . . . .	14
<b>11 Implementation Roadmap</b>	<b>14</b>
11.1 MVP (0–3 months) . . . . .	14
11.2 Phase 2 (3–6 months) . . . . .	14
11.3 Phase 3 (6–12 months) . . . . .	14
11.4 Enterprise Rollout (12+ months) . . . . .	14
<b>12 Free-Tier Friendly Notes</b>	<b>15</b>
<b>13 Database &amp; Offline-Edge Strategy</b>	<b>15</b>
13.1 MongoDB Atlas (Central Database) . . . . .	15
13.2 SQLite (Local / Edge Cache) . . . . .	15
13.3 Combined Workflow . . . . .	16
<b>14 Satellite Data Alternatives &amp; Local Simulation (Sentinel-2)</b>	<b>16</b>
14.1 Options for Free Sentinel-2 Data Without Google Cloud Registration . . . . .	16
14.2 Local Simulation for MVP . . . . .	17
14.3 Recommended MVP Strategy . . . . .	17
<b>15 APIs and External Integrations</b>	<b>17</b>
15.1 Core APIs . . . . .	17
<b>16 All Other Requirements</b>	<b>18</b>
16.1 Libraries and Dependencies . . . . .	18
16.2 Hardware and Infrastructure . . . . .	18
16.3 Testing and Quality Assurance . . . . .	18
16.4 Additional Considerations . . . . .	19
<b>17 Summary</b>	<b>19</b>

## Introduction

The AI Crop Monitoring & Agricultural Intelligence Platform is a professional-grade ecosystem for precision agriculture. It integrates computer vision, time-series forecasting, market analytics, explainable AI, and microfinance scoring into a single, modular, and scalable system.

It is designed to serve:

- Farmers: simple mobile/web interface.
- Cooperatives/NGOs: dashboards for group management.
- Enterprises: large-scale deployments with monitoring, compliance, and integration into supply chains.

The system combines edge computing (offline models on phones, Raspberry Pi, or drones) with cloud-based retraining (Google Colab, Kaggle, or free-tier cloud). This makes it both low-cost and scalable.

## Project Structure

You will maintain the following directory structure (already adapted for professional projects):

```
C:\Users\ADEGOKE\Desktop\Crop Monitor
|-- LICENSE
|-- README.md
|-- Dockerfile
|-- docker-compose.yml
|-- docker-compose.prod.yml
|-- docker-compose.mlflow.yml
|-- pyproject.toml
|-- requirements.txt
|-- run_pipeline.py
|-- setup.py
|-- main.py
|-- .gitignore
|-- .github/workflows/ci-cd.yml
|
|-- notebooks/ (experiments, EDA, training)
|-- monitoring/ (Prometheus, Grafana configs)
|-- docs/ (guides, architecture, API docs)
|-- deployment/ (Terraform, Kubernetes, configs)
|-- infrastructure/ (Docker, Helm, Terraform, manifests)
|-- src/ (recommendation_engine, frontend, crop_monitor)
|-- backend/ (workers, APIs)
|-- ai/ (training, models, explainability)
|-- tests/ (unit, integration, performance)
```

## When & How You'll Use It

- Use notebooks/ for experiments in Google Colab or Jupyter. This includes exploratory data analysis (EDA), model prototyping, and initial training scripts.
- Use src/ for production-ready code. This directory houses modular components like recommendation engines and core business logic.
- Use ai/ for training pipelines and model definitions. Includes scripts for model training, evaluation, and export to formats like ONNX or TFLite.
- Use docs/ to store guides for farmers, administrators, and deployment engineers. Utilize Sphinx or MkDocs for generating documentation from Markdown files.
- Use tests/ to ensure everything works before deployment. Implement unit tests with pytest, integration tests for API endpoints, and performance tests for model inference times.

## Features & Capabilities

Category	Feature	Why It Matters	When & How You Use It
Data Sources	Multi-modal (images, soil, weather, satellite, drone)	Improves prediction accuracy	Download from Kaggle (PlantVillage/PlantDoc), FAO SoilGrids, Sentinel-2 (Google Earth Engine or alternatives), OpenWeatherMap API. Use DVC to track versions and ensure reproducibility.
Camera Inputs	Smartphone, webcam, drone, IP camera	Enables scalable real-time capture	Farmers upload via Streamlit/Next.js UI. Developers test with OpenCV and webcams. Process images using cv2.imread for local testing.
Explainable AI	SHAP, LIME, Grad-CAM	Builds trust with farmers	Use Grad-CAM in notebooks to show heatmaps; use SHAP on structured models (e.g., credit risk). Integrate into dashboards for visual explanations.
Recommendations	Pesticide, fertilizer, irrigation	Provides actionable insights	Define mapping in YAML/JSON configs (src/config/treatments.yaml). Load configs using PyYAML and generate recommendations based on model outputs.

Continuous Learning	Feedback loops + active learning	Models evolve over time	Collect farmer feedback into SQLite (offline) and sync to MongoDB Atlas; retrain models monthly using Colab free GPUs. Use active learning strategies like uncertainty sampling.
Offline Mode	Edge inference	Enables remote farm usage	Export models to TensorFlow Lite / ONNX and deploy on mobile/edge devices. Use tflite_runtime for inference on Android/iOS, store results in SQLite.
Predictive Analytics	Yield forecasting, disease spread	Improves crop planning	Use Prophet or LSTM in ai/models/yield_forecasting/. Train in Colab, deploy via FastAPI endpoints for API access, store results in MongoDB Atlas.
Market Insights	Dynamic pricing, supply chain	Maximizes farmer profit	Collect price data via scraping (BeautifulSoup/Selenium); apply regression models like LinearRegression from scikit-learn. Store in MongoDB Atlas market_prices collection. Schedule with Celery.
Digital Finance	Loan recommendations	Provides financial access	Train logistic regression or boosting models; deploy via backend/api/recommendations.py. Store in MongoDB Atlas loan_requests collection.
Notifications	WhatsApp, Email, SMS	Ensures timely actions	Use Twilio sandbox for SMS/WhatsApp and Gmail API for email. Store logs in MongoDB Atlas notifications collection. Integrate with Celery for asynchronous sending.
Scalability	Multi-farm dashboards	Enterprise readiness	Build dashboards in Streamlit (src/frontend/streamlit_app/dashboard.py). Fetch data from MongoDB Atlas for multi-farm views. Use Kubernetes for scaling.

Compliance	GDPR-style logging	Trust & legal safety	Encrypt farmer logs in MongoDB Atlas crop_monitoring collection; allow farmers to request deletion. Use PyMongo for data access and implement audit logs.
------------	--------------------	----------------------	---

## Next-Level Differentiators

- **Edge AI & On-Farm Inference:** Deploy compressed models (ONNX, TFLite) on smartphones, Raspberry Pi, or drones. Use ONNX Runtime for cross-platform inference.
- **Adaptive Treatment Optimization:** Store historical actions in MongoDB Atlas → retrain dosage optimization model with new outcomes using reinforcement learning techniques if advanced.
- **AI Pest/Disease Forecasting:** Use weather API + LSTM model to forecast outbreaks. Implement in ai/models/forecasting.py with time-series data stored in MongoDB Atlas yield\_forecasts collection.
- **Satellite + Drone Fusion:** Process Sentinel-2 via Google Earth Engine or alternatives; fuse with drone images in fusion\_pipeline.py using libraries like rasterio for alignment. Store metadata in MongoDB Atlas images\_metadata collection.
- **Farm Health Score:** Weighted index (disease % + soil + rainfall). Compute in back-end/services/health\_score.py and store in MongoDB Atlas farms collection for dashboard display.
- **AI Marketplace Integration:** Web scrape suppliers using Selenium and store in MongoDB Atlas market\_prices collection for integration into recommendations.
- **Gamification & Community:** Reward farmers for consistent use (Streamlit/Gradio leaderboards). Track points in MongoDB Atlas farmer\_feedback collection and display top users.
- **Climate Analytics:** Show sustainability metrics (CO<sub>2</sub>, fertilizer usage) using open APIs like Carbon Interface API. Store results in MongoDB Atlas for analytics.
- **Voice Interfaces:** Use Whisper (offline STT) or Vosk for local languages (Yoruba, Hausa, Igbo). Store voice command logs in MongoDB Atlas for analysis.
- **AR Overlays:** Use AR.js/ARCore to guide farmers visually on infected leaf regions. Store AR metadata in MongoDB Atlas images\_metadata collection.
- **CI/CD for Models:** GitHub Actions checks accuracy before pushing models live. Store model metadata in MongoDB Atlas for tracking.

## Governance, Compliance & Reliability

### Licensing & Provenance

When: Every time you add or update datasets.

How:

- Maintain manifest.csv with source, license, date, sha256 checksum. Use pandas to generate and validate, store metadata in MongoDB Atlas crop\_monitoring collection.
- Use DVC for dataset versioning. Commands: dvc add data/, dvc push.
- Require farmer consent forms for uploads. Store digitally in MongoDB Atlas users collection with timestamps.

## Annotation & Label QA

When: Before training any supervised model.

How:

- Use CVAT or Roboflow for annotations. Export in COCO or YOLO formats, store in MongoDB Atlas images\_metadata collection.
- Set clear guidelines: resolution (min 1024x1024), label taxonomy (e.g., disease classes from PlantVillage).
- Ensure inter-annotator agreement 0.75 (Cohen's Kappa). Compute using scikit-learn's cohen\_kappa\_score.

## Data Validation

When: Before training and deployment.

How:

- Add Great Expectations tests: correct file counts, schema checks. Define expectations in great\_expectations/expectations/. Validate MongoDB Atlas data.
- Run pytest in CI/CD pipelines. Integrate with GitHub Actions for automated validation.

## Model Governance & Cards

When: At each model release.

How:

- Store model card (model-card.md) with intended use, limitations, metrics (accuracy, F1-score). Save in MongoDB Atlas crop\_monitoring collection.
- Track in MLflow registry. Use mlflow.log\_artifact for cards, sync metadata to MongoDB Atlas for multi-user access.

## Drift Detection & Retraining

When: Continuously after deployment.

How:

- Use Evidently AI to detect distribution drifts. Generate reports in monitoring/evidently/, store in MongoDB Atlas system\_logs collection.
- Retrain if macro-F1 drops below baseline. Automate with cron jobs or Airflow, update models in MongoDB Atlas.

## **Security, Privacy & Compliance**

When: Always in production.

How:

- Store secrets in Vault/GCP Secret Manager. Access via API in code.
- Encrypt data (AES-256 at rest, TLS in transit). Use MongoDB Atlas client-side field-level encryption.
- Implement data retention policies (max 2 years). Use MongoDB Atlas queries to purge old data.

## **Observability & SLOs**

When: During production monitoring.

How:

- Set latency p95 <200ms, uptime 99.5%.
- Monitor with Prometheus, Grafana, OpenTelemetry. Store metrics in MongoDB Atlas system\_logs collection.

## **CI/CD & Testing Gates**

When: Before merging to production.

How:

- GitHub Actions → Run tests (pytest, lint with black/flake8).
- Block deployment if model accuracy < baseline. Use conditional steps in workflows.
- Canary rollout with rollback triggers. Implement using Kubernetes deployments.

## **Reproducibility & Backups**

When: Every training run & release.

How:

- Pin seeds, dependencies (requirements.txt, pyproject.toml). Use random.seed(42) in code.
- Version datasets with hashes. Use DVC or Git LFS, store metadata in MongoDB Atlas.
- Backup MongoDB Atlas weekly (RPO=1 day, RTO=4h). Use mongodump for exports, automated via cron.

## **Knowledge & Skills Roadmap**

### **Mathematics & Statistics**

- Linear Algebra, Calculus → For CNN/LSTM backpropagation. (Resource: Gilbert Strang, MIT OCW lectures on YouTube or website.)
- Probability & Statistics → For risk scoring & uncertainty. (Resource: "Introduction to Statistical Learning" by James et al., free PDF available.)

- Information Theory → For model optimization (entropy, KL). (Resource: "Elements of Information Theory" by Cover & Thomas.)

## Machine Learning

- Trees, Random Forests → Market pricing, yield risk. (scikit-learn docs: <https://scikit-learn.org/stable/modules/ensemble.html>)
- Boosting (XGBoost, LightGBM) → Credit scoring, yield models. (Kaggle tutorials: Search for XGBoost on Kaggle kernels.)
- Clustering → Group farms by conditions. (Use KMeans from scikit-learn.)

## Deep Learning

- CNNs (ResNet, EfficientNet) → Crop disease detection. (PyTorch/TensorFlow tutorials: Official docs on transfer learning.)
- Transformers (ViT) → High-performance image classification. (Hugging Face Transformers library.)
- LSTM/GRU → Weather & yield forecasting. (Keras sequential models.)

## MLOps

- MLflow → Experiment tracking. (Docs: <https://mlflow.org/docs/latest/index.html>)
- Docker/Kubernetes → Deployment. (Official tutorials for containerization.)
- Prometheus/Grafana → Monitoring. (Setup guides on official sites.)

## Edge AI

- ONNX, TFLite → Compress & deploy on phones. (Conversion scripts in TensorFlow/PyTorch.)
- Pruning, Quantization → Reduce model size. (Use `torch.nn.utils.prune`.)

## Compliance

- GDPR basics → Data privacy. (Read EU GDPR documentation.)
- OAuth2/JWT → Secure farmer login. (Use Authlib or FastAPI Users.)

## Frontend Details

The frontend provides user-friendly interfaces for different stakeholders. It is built to be responsive, accessible, and integrated with backend APIs.

## Technologies and Frameworks

- **Streamlit:** For rapid development of dashboards and prototypes. Ideal for data scientists to create interactive apps quickly.

- **Next.js with React:** For production web applications. Supports server-side rendering, API routes, and static generation for performance.
- **React Native:** For cross-platform mobile apps (iOS/Android). Enables native camera access and offline capabilities.
- **UI Components:** Material-UI (MUI) for consistent design, or Tailwind CSS for custom styling.
- **Visualization Libraries:** Chart.js for charts, Leaflet.js or Mapbox for maps displaying farm locations and satellite overlays.
- **State Management:** Redux or MobX for complex states in Next.js/React Native.
- **AR and Voice:** ARCore/ARKit for mobile AR, Web Speech API or Vosk.js for voice in web.

## Key Components and Pages

- **Login/Register:** OAuth2 integration with email/phone verification.
- **Farmer Dashboard:** Image upload form, real-time recommendations, notifications feed, farm health score visualization.
- **Cooperative Management:** Multi-user views, aggregated analytics, group leaderboards for gamification.
- **Enterprise Console:** Compliance reports, scalability metrics, integration settings for supply chains.
- **Mobile Specific:** Camera integration for direct captures, offline mode with local storage (SQLite via Expo SQLite), sync when online to MongoDB Atlas.

## Code Structure

Located in src/frontend/.

- streamlit\_app/: dashboard.py for main app, pages/ for multi-page setup.
- nextjs\_app/: pages/ for routes (e.g., index.js for home), components/ for reusable UI (e.g., ImageUploader.js).
- reactnative\_app/: App.js entry, screens/ for views (e.g., DashboardScreen.js), services/ for API calls.

## Example Code Snippet (Streamlit Upload)

```
import streamlit as st
import requests

st.title("CropedImageUpload")
uploaded_file = st.file_uploader("Choose an image...", type=["jpg", "png"])
if uploaded_file is not None:
    response = requests.post("http://backend/api/inference", files={"file": uploaded_file})
    st.write(response.json()["recommendation"])
```

## Deployment and Testing

- Deploy Streamlit on Heroku or AWS; Next.js on Vercel; React Native via Expo or direct builds.
- Testing: Jest for unit tests, Cypress for E2E tests.
- Accessibility: Ensure WCAG compliance with ARIA labels.

## Backend Architecture (Atlas Integrated)

The backend manages business logic, API endpoints, and background processing, with MongoDB Atlas as the primary database.

## Technologies and Frameworks

- **FastAPI:** Asynchronous API framework for high performance.
- **Celery:** Task queue for asynchronous jobs like model training or scraping.
- **Database Connections:** MongoDB Atlas via pymongo for centralized storage.
- **Workers:** Dedicated scripts for compute-intensive tasks (e.g., image processing with OpenCV).
- **Authentication:** FastAPI Security with JWT, dependencies for role-based access (farmer, admin).

## MongoDB Atlas Collections

- **crop\_monitoring**
  - images\_metadata: Stores image paths, labels, GPS, drone info.
  - farmer\_feedback: User feedback for retraining and recommendation refinement.
  - market\_prices: Scrapped or API-based market data.
  - loan\_requests: Microfinance and credit scoring data.
  - notifications: SMS/email/WhatsApp logs.
  - yield\_forecasts: Predicted crop yields.
  - system\_logs: Monitoring and drift detection logs.

## Use Cases

- Query images /  
*metadata for dashboard display. Collect farmer feedback in real-time for active learning.*
- Update market\_prices for dynamic pricing models. Track loan\_request history and repayment.
- Monitor recommendation usage via notifications.

## Key Endpoints and Services

- **/api/auth:** Login, register, token refresh.

- **/api/upload**: Image upload, trigger inference, store metadata in MongoDB Atlas images\_metadata.
- **/api/recommendations**: Fetch personalized suggestions from MongoDB Atlas based on user ID.
- **/api/forecast**: POST with farm data, store results in MongoDB Atlas yield\_forecasts.
- **/api/notifications**: Send alerts, log in MongoDB Atlas notifications collection.
- **/api/market**: Get current prices from MongoDB Atlas market\_prices, historical trends.
- **Background Services**: Scraper worker for market data, inference worker for models.

## Code Structure

Located in backend/.

- api/: main.py for app, routers/ for endpoints (e.g., inference.py).
- workers/: tasks.py for Celery tasks.
- services/: Business logic (e.g., recommendation\_service.py).
- models/: Pydantic models for request/response validation.

## Example Code Snippet (FastAPI with MongoDB Atlas)

```
from fastapi import FastAPI, UploadFile, File
from pymongo import MongoClient
import tensorflow as tf

app = FastAPI()
client = MongoClient("mongodb+srv://<user>:<pass>@cluster0.mongodb.net/")
db = client.crop_monitoring

@app.post("/api/inference")
async def inference(file: UploadFile = File(...)):
    content = await file.read()
    img = tf.image.decode_image(content)
    prediction = model.predict(img)
    result = {"disease": prediction.argmax(), "confidence": float(
        prediction.max())}
    db.images_metadata.insert_one({"image": file.filename, "result": result})
    return result
```

## Deployment and Testing

- Run locally with uvicorn backend.api.main:app --reload.
- Scale with Gunicorn, deploy in Docker containers.
- Testing: Pytest for APIs, Locust for load testing.

## Edge & Offline Devices

This section details the strategy for edge computing and offline functionality, ensuring usability in remote areas.

### Edge Models

- **TFLite / ONNX:** Deploy compressed models on smartphones, Raspberry Pi, or drones for local inference.
- **Implementation:** Use `tflite_runtime` or ONNX Runtime for inference. Convert models in notebooks/`export.py`.

### Caching

- **SQLite:** Local storage on edge devices for offline data (images, predictions, recommendations).
- **JSON:** Alternative lightweight format for simple metadata if SQLite is too heavy.
- **Sync to MongoDB Atlas:** When network is available, push SQLite/JSON data to Atlas collections (`images_metadata`, `farmer_feedback`).

### Offline Recommendations

- Deliver immediate recommendations using local TFLite/ONNX models.
- Store results in SQLite, sync to MongoDB Atlas when online.

### Workflow Example

- Farmer captures image → save to SQLite with metadata.
- Run inference locally → store result in SQLite.
- On network availability → sync SQLite to MongoDB Atlas, clear synced rows.

## Observability & CI/CD

This section ensures robust monitoring and deployment pipelines.

### Monitoring

- **Prometheus + Grafana:** Track API latency, model performance, and system health. Store metrics in MongoDB Atlas `system_logs`.
- **OpenTelemetry:** For distributed tracing across services.

### Model Tracking

- **MLflow:** Track experiments locally, sync metadata to MongoDB Atlas `crop_monitoring` for multi-user setups.
- **Model Cards:** Store in MongoDB Atlas with metrics and limitations.

## CI/CD

- **GitHub Actions:** Run tests, linting, and model accuracy checks before deployment.
- **Workflows:** Defined in .github/workflows/ci-cd.yml. Block if accuracy < baseline.
- **Canary Rollouts:** Use Kubernetes for staged deployments with rollback triggers.

## Alerts & Logs

- Store alerts in MongoDB Atlas notifications collection.
- System logs (errors, drifts) in MongoDB Atlas system\_logs collection.

## Implementation Roadmap

### MVP (0–3 months)

- Dataset setup (PlantVillage). Download from Kaggle, organize in data/, store metadata in MongoDB Atlas.
- Train CNN (ResNet50). Use transfer learning in notebooks/train\_cnn.ipynb.
- Build Streamlit app for farmers. Implement upload and inference in streamlit\_app.py, query MongoDB Atlas.
- Feedback collection in SQLite (offline), sync to MongoDB Atlas farmer\_feedback.

### Phase 2 (3–6 months)

- Add yield forecasting (Prophet/LSTM). Script in ai/yield\_forecast.py, store in MongoDB Atlas yield\_forecasts.
- Add market pricing model. Scraper in backend/scrapers/market.py, store in MongoDB Atlas market\_prices.
- Add microfinance scoring (logistic regression). Model in ai/finance/, store in MongoDB Atlas loan\_requests.

### Phase 3 (6–12 months)

- Integrate satellite + drone pipelines. Fusion script with GDAL/rasterio, store in MongoDB Atlas images\_metadata.
- Launch cooperative dashboards. Advanced Streamlit or Next.js views fetching from MongoDB Atlas.
- Add gamification, voice interfaces. Leaderboards in frontend, Vosk integration, logs in MongoDB Atlas.

### Enterprise Rollout (12+ months)

- Deploy full CI/CD. Complete workflows in .github/.
- Implement drift detection + retraining automation. Store metrics in MongoDB Atlas system\_logs.

- Launch with NGOs & government programs. Prepare demo and documentation, all data centralized in MongoDB Atlas.

## Free-Tier Friendly Notes

- **MongoDB Atlas Free-Tier (M0)**: 512MB storage, sufficient for MVP data and multi-farm prototype. Scales automatically for small datasets.
- **Google Colab / Kaggle**: Free GPUs for model training. Use notebooks/ for experiments.
- **Edge Devices + SQLite**: Offline-first with local storage, sync to MongoDB Atlas when online.
- **Cost Efficiency**: No payment or ATM verification needed for MVP; leverages free tiers of AWS, MongoDB Atlas, and open APIs.

## Database & Offline-Edge Strategy

### MongoDB Atlas (Central Database)

**Purpose:** Centralized cloud database for multi-user access, dashboards, analytics, and persistent storage.

#### Collections:

- images\_metadata: Stores image paths, labels, GPS, drone info.
- farmer\_feedback: User feedback for retraining and recommendation refinement.
- market\_prices: Scrapped or API-based market data.
- loan\_requests: Microfinance and credit scoring data.
- yield\_forecasts: Predicted crop yields.
- notifications: SMS/email/WhatsApp logs.
- system\_logs: Monitoring, drift detection, and system metrics.

**Usage:** All multi-user, cross-device, or long-term analytics stored in MongoDB Atlas. Access via pymongo with connection string in .env.

### SQLite (Local / Edge Cache)

**Purpose:** Temporary, local storage on mobile/edge devices (smartphones, Raspberry Pi, drones).

**Why:** Supports offline-first operations; avoids blocking farm use when there is no internet.

#### Example Data:

- Uploaded images before syncing to MongoDB Atlas.
- Local predictions and recommendations.
- Sensor readings or drone metadata.

#### Edge Workflow:

- Farmer captures image offline → save metadata in SQLite.
- Device computes local inference using TFLite/ONNX model → store in SQLite.
- When internet is available → sync SQLite entries to MongoDB Atlas (images\_metadata, farmer\_feedback).
- Clean SQLite after successful sync to manage storage.

**Notes:** MongoDB Atlas is the single source of truth; SQLite is a temporary offline buffer.

## Combined Workflow

- **Edge Device / Offline Mode:**
  - Image + Metadata → SQLite.
  - Local TFLite/ONNX inference.
  - Recommendation / Feedback → SQLite.
- **Sync Process (when online):**
  - SQLite rows → MongoDB Atlas.
  - Confirm success → Clear synced SQLite rows.
  - Dashboards / Backend read from MongoDB Atlas.

### Benefits:

- Ensures offline-first usability for farmers.
- Maintains centralized data in MongoDB Atlas for analytics, dashboards, and retraining.
- Low-cost, free-tier compatible with MongoDB Atlas M0.

## Satellite Data Alternatives & Local Simulation (Sentinel-2)

### Options for Free Sentinel-2 Data Without Google Cloud Registration

Since Google Earth Engine registration is not possible, use these alternatives:

- **AWS Open Data ([registry.opendata.aws/sentinel-2/](https://registry.opendata.aws/sentinel-2/)):** Public S3 buckets for Sentinel-2 data. Access via boto3: `s3 = boto3.client('s3');` `s3.download_file('sentinel-s2-l1c', path, local).`
- **Copernicus Open Access Hub ([scihub.copernicus.eu](https://scihub.copernicus.eu)):** Official ESA portal for free Sentinel downloads. Register (free), use API with `sentinelsat`: from `sentinelsat import SentinelAPI;` `api.query(...);` `api.download()`.
- **USGS EarthExplorer ([earthexplorer.usgs.gov](https://earthexplorer.usgs.gov)):** Free access to Sentinel-2 and Landsat datasets. Bulk download tools available.
- **Sentinel Hub ([www.sentinel-hub.com](https://www.sentinel-hub.com)):** Free tier for browsing, analyzing, and downloading Sentinel-2 data. Use EO Browser or `sentinelhub-py` library.
- **SkyWatch EXPLORE:** Aggregates multiple sources, free for certain datasets.

**Notes:** These sources provide .SAFE or GeoTIFF formats without requiring cloud payments.

## Terrestrial Data Alternatives

### Local Simulation for MVP

If live Sentinel-2 access is unavailable:

- **Raster Images:**
  - Use sample .tif or .png bands from AWS/Copernicus datasets.
  - Store in data/satellite/ for local processing.
- **CSV Feature Extraction:**
  - Precompute indices like NDVI, EVI using rasterio or GDAL.
  - Save as satellite\_features.csv for ML model input.
- **Edge Device Testing:**
  - Use sample CSV + TFLite model for offline yield prediction or disease risk scoring.
  - Ensures functionality without cloud connectivity.

### Recommended MVP Strategy

- Start offline-first using SQLite + sample Sentinel-2 rasters or CSVs.
- Sync all offline outputs to MongoDB Atlas when network is available.
- Replace simulation with real satellite feeds when Google Earth Engine or other APIs become accessible.

## APIs and External Integrations

All APIs and services required for data ingestion and notifications.

### Core APIs

- **OpenWeatherMap API:** For real-time and forecast weather. Sign up for free API key, limit to 1,000 calls/day. Usage: `requests.get(f"https://api.openweathermap.org/data/2.5/weather?lat={lat}&lon={lon}&appid={key}")`. Store in MongoDB Atlas yield\_forecasts.
- **Twilio API:** SMS/WhatsApp. Free sandbox for testing, paid for production. Client: `from twilio.rest import Client; client.messages.create(...)`. Log in MongoDB Atlas notifications.
- **Gmail API:** Email notifications. Use OAuth2 with google-auth library, store logs in MongoDB Atlas.
- **BeautifulSoup/Selenium:** For scraping market prices from agricultural sites. Store in MongoDB Atlas market\_prices.
- **FAO APIs:** For soil data (SoilGrids). Free access via REST endpoints, store in MongoDB Atlas.
- **PubChem API:** For chemical info on pesticides/fertilizers, integrate into recommendations.

# All Other Requirements

## Libraries and Dependencies

Full requirements.txt:

```
fastapi==0.95.0
uvicorn==0.21.0
streamlit==1.20.0
tensorflow==2.11.0
torch==1.13.0
scikit-learn==1.2.0
xgboost==1.7.0
lightgbm==3.3.0
prophet==1.1.0
opencv-python==4.7.0
shap==0.40.0
lime==0.2.0
mlflow==2.1.0
dvc==2.45.0
great-expectations==0.15.0
evidently==0.2.0
prometheus-client==0.16.0
pymongo==4.3.0
celery==5.2.0
twilio==7.16.0
beautifulsoup4==4.11.0
selenium==4.8.0
rasterio==1.3.0
geopandas==0.12.0
onnx==1.13.0
tensorflow-lite-runtime==2.11.0
vosk==0.3.0
requests==2.28.0
pandas==1.5.0
numpy==1.24.0
matplotlib==3.7.0
sentinelhub==3.9.0
sentinelsat==1.1.0
boto3==1.26.0
```

## Hardware and Infrastructure

- **Edge Devices:** Raspberry Pi 4 (for on-farm inference), drones like DJI with cameras.
- **Cloud:** AWS Free Tier (EC2 t2.micro, S3), MongoDB Atlas M0, alternatives like DigitalOcean.
- **Tools:** Git for version control, Docker for containerization, Kubernetes/Helm for orchestration, Terraform for IaC.

## Testing and Quality Assurance

- Unit Tests: Pytest for functions, models.

- Integration: Test API chains with MongoDB Atlas.
- Performance: Measure inference time, scalability.
- Security: OWASP scans, dependency checks with pip-audit.

## Additional Considerations

- **Cost Management:** Use free tiers of MongoDB Atlas, AWS, and open APIs.
- **Localization:** Support Yoruba, Hausa, Igbo in UI (React Intl).
- **Sustainability:** Optimize models for low energy on edge devices.

## Summary

- **Training Pipeline:** Local images + CSV manifests, processed in Colab/Kaggle.
- **Backend & Multi-User Functionality:** MongoDB Atlas for centralized storage and analytics.
- **Edge & Offline:** TFLite/ONNX models + SQLite, synced to MongoDB Atlas.
- **Dashboards & Analytics:** FastAPI/Streamlit reading from MongoDB Atlas.
- **MVP Compatibility:** Fully free-tier compatible without payment or ATM verification.